



IPv6 Programming Introduction


Linjiang Li

6/19 2014

IPv6 and why is IPv6?

- **Internet Protocol version 6 (IPv6)** , which is intended to replace IPv4.
- With more and more new devices being connected to the Internet, the IPv4 address will be exhaustion.
- IPv6 uses a 128-bit address, allowing 2^{128} , or approximately 3.4×10^{38} addresses, or more than 7.9×10^{28} times as many as IPv4, which uses 32-bit addresses. IPv4 provides approximately 4.3 billion addresses.
- The number of IPv6 addresses would be about 40,000 addresses for every atom on the surface of the earth.

IPv6 Addressing



IPv6 Addressing

IPv4 32-bits

IPv6 128-bits


$$2^{32} = 4,294,967,296$$

$$2^{128} = 340,282,366,920,938,463,463,374,607,431,768,211,456$$

$$2^{128} = 2^{32} * 2^{96}$$

$$2^{96} = 79,228,162,514,264,337,593,543,950,336 \text{ times the number of possible IPv4 Addresses (79 trillion trillion)}$$

IPv6 Addressing



World's population is
approximately 6.5 billion

$$\frac{2^{128}}{6.5 \text{ Billion}} = 52 \text{ Trillion Trillion IPv6 addresses per person}$$



Typical brain has
~100 billion brain cells
(your count may vary)

$$\frac{52 \text{ Trillion Trillion}}{100 \text{ Billion}} = 523 \text{ Quadrillion (523 thousand trillion) IPv6 addresses for every human brain cell on the planet!}$$

IPv6 Address Format

Representation

- 16-bit hexadecimal numbers
- Numbers are separated by (:)
- Hex numbers are not case sensitive
- Abbreviations are possible

Leading zeros in contiguous block could be represented by (::)

Example:

2001:0db8:0000:130F:0000:0000:087C:140B

2001:0db8:0:130F::87C:140B

Double colon only appears once in the address

IPv6 Address Prefix

Prefix Representation

- Representation of prefix is just like CIDR
- In this representation you attach the prefix length
- Like v4 address:

198.10.0.0/16

- V6 address is represented the same way:

2001:db8:12::/48

- Only leading zeros are omitted. Trailing zeros are not omitted

2001:0db8:0012::/48 = 2001:db8:12::/48

2001:db8:1200::/48 ≠ 2001:db8:12::/48

IPv6 Address Model

- Addresses are assigned to interfaces

Change from IPv4 mode:

- Interface “expected” to have multiple addresses

- Addresses have scope

Link Local

Unique Local

Global

- Addresses have lifetime

Valid and preferred lifetime



IPv6 Address Types

- Unicast

Address of a single interface. One-to-one delivery to single interface

- Multicast

Address of a set of interfaces. One-to-many delivery to all interfaces in the set

- Anycast

Address of a set of interfaces. One-to-one-of-many delivery to a single interface in the set that is closest

- No more broadcast addresses

Some Special Addresses

Symbolic Name (IPv4 / IPv6)	IPv4 Address	IPv6 Address
INADDR_ANY / IN6ADDR_ANY_INIT	0.0.0.0	::
INADDR_LOOPBACK / IN6ADDR_LOOPBACK_INIT	127.0.0.1	::1
INADDR_BROADCAST	255.255.255.255	Not exist

➤ **More IPv4/IPv6 information, refer attached.**

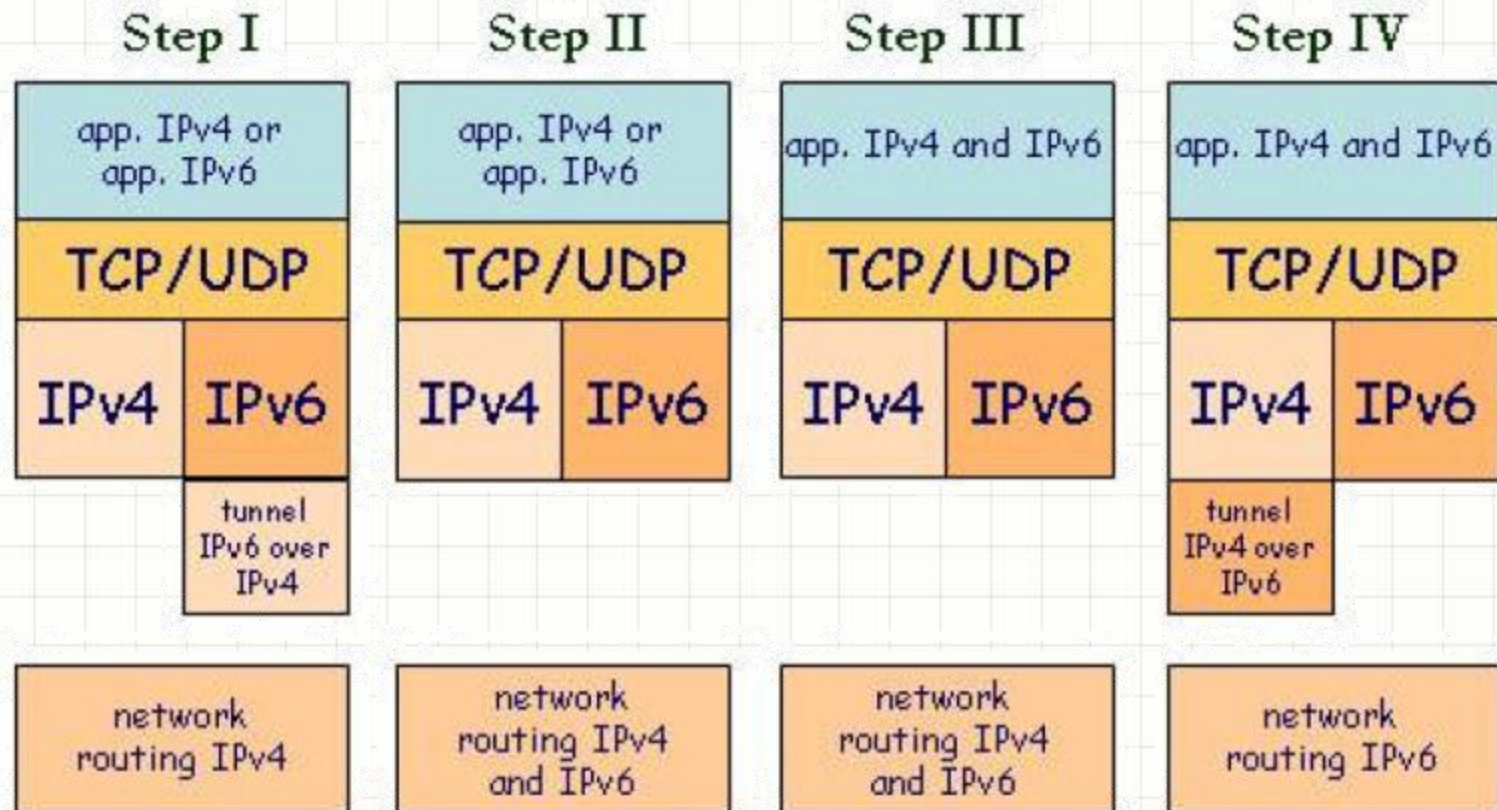


ipv6_reference_ca
rd.pdf

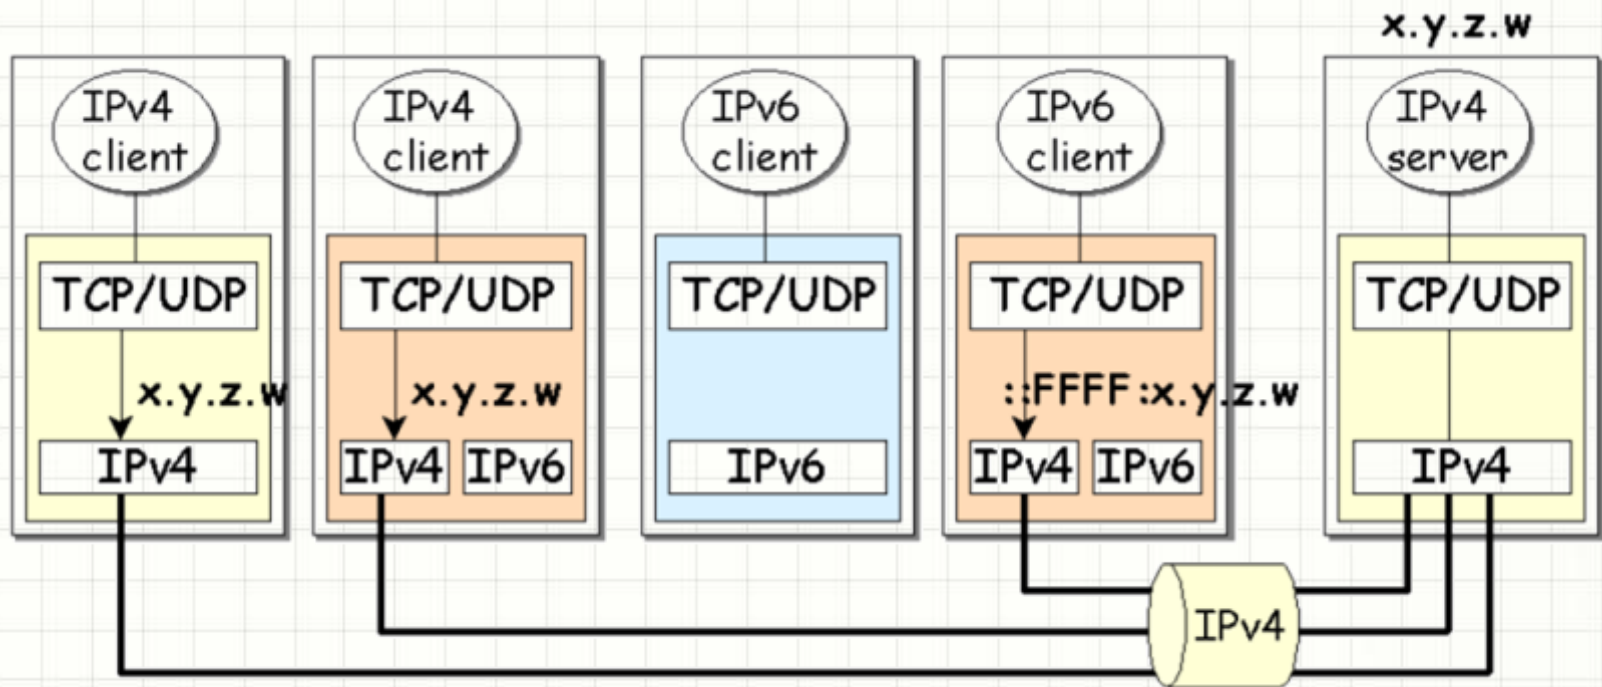


Transition & Interoperability

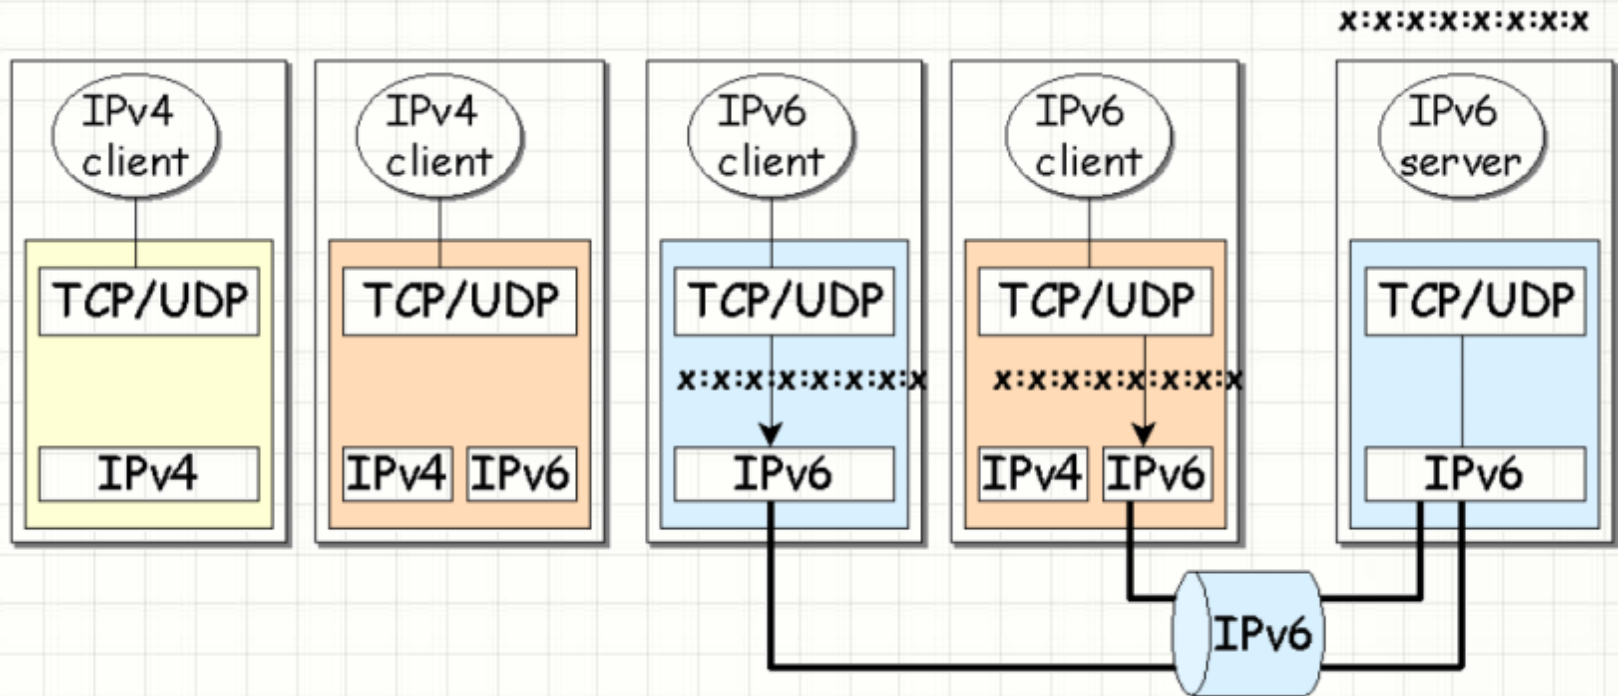
Transition scenario



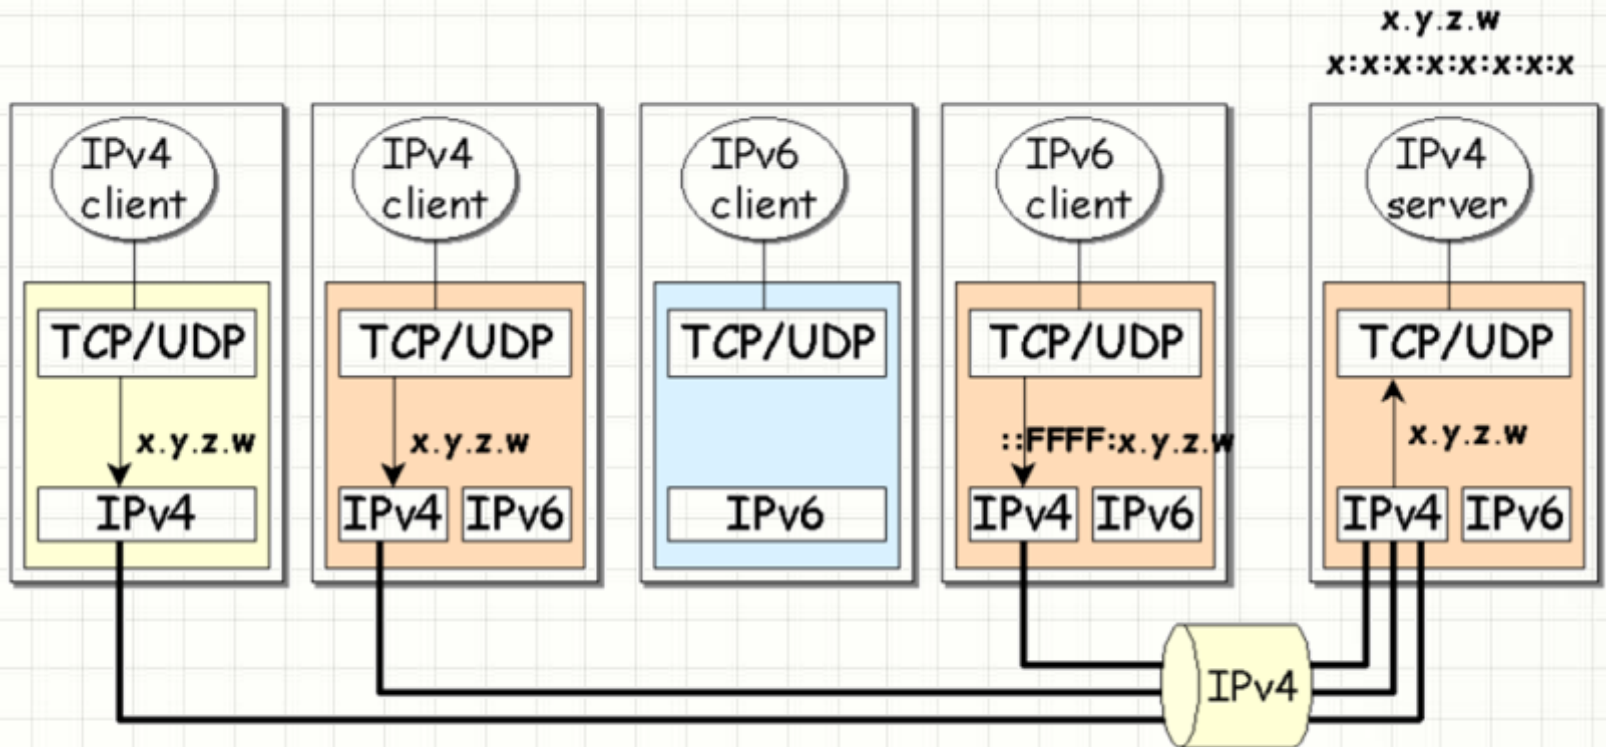
IPv4/IPv6 Interoperability (IPv4 server at IPv4 only node)



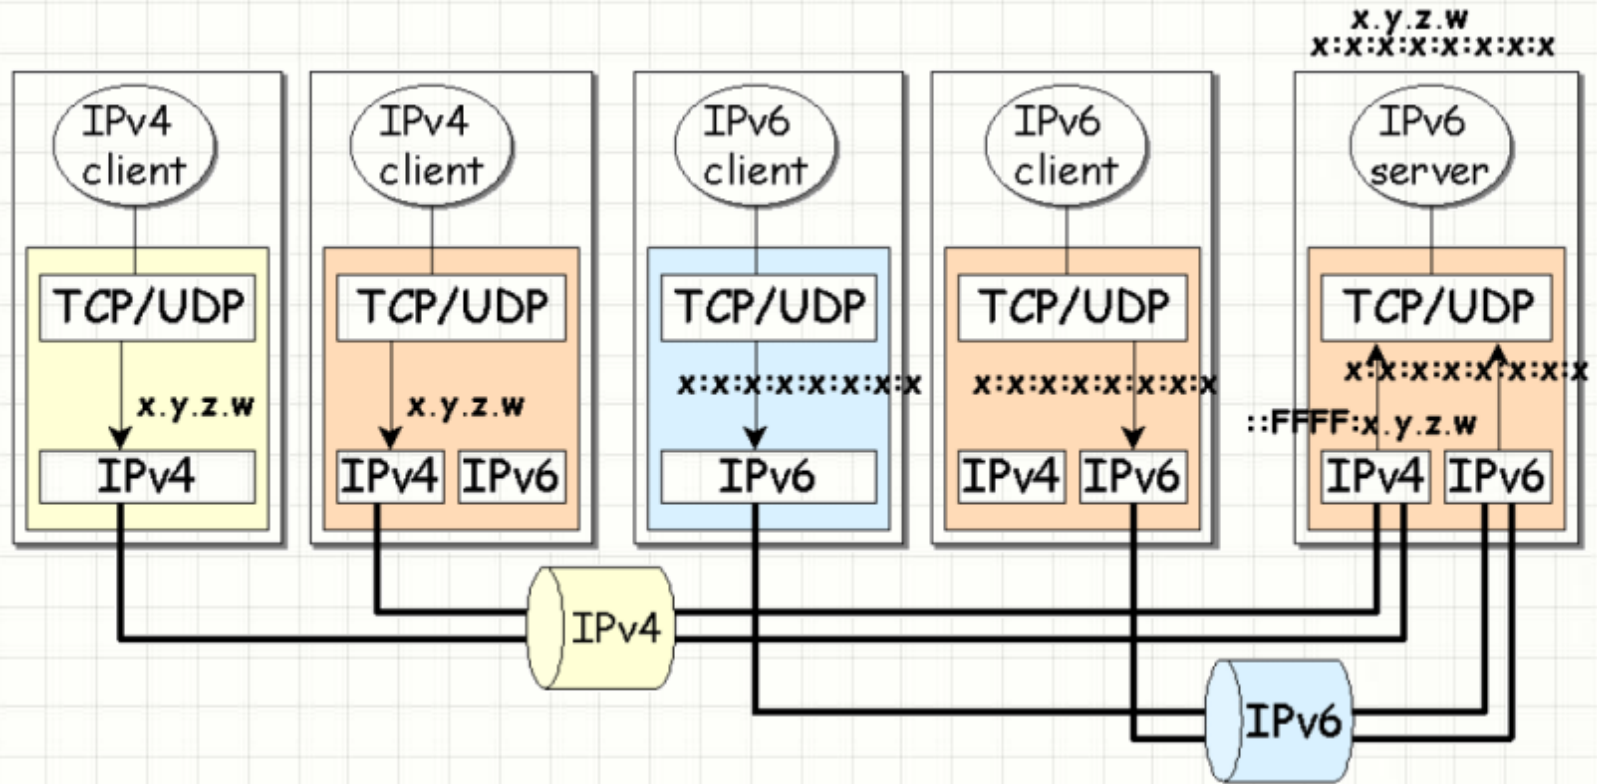
IPv4/IPv6 Interoperability (IPv6 server at IPv6 only node)



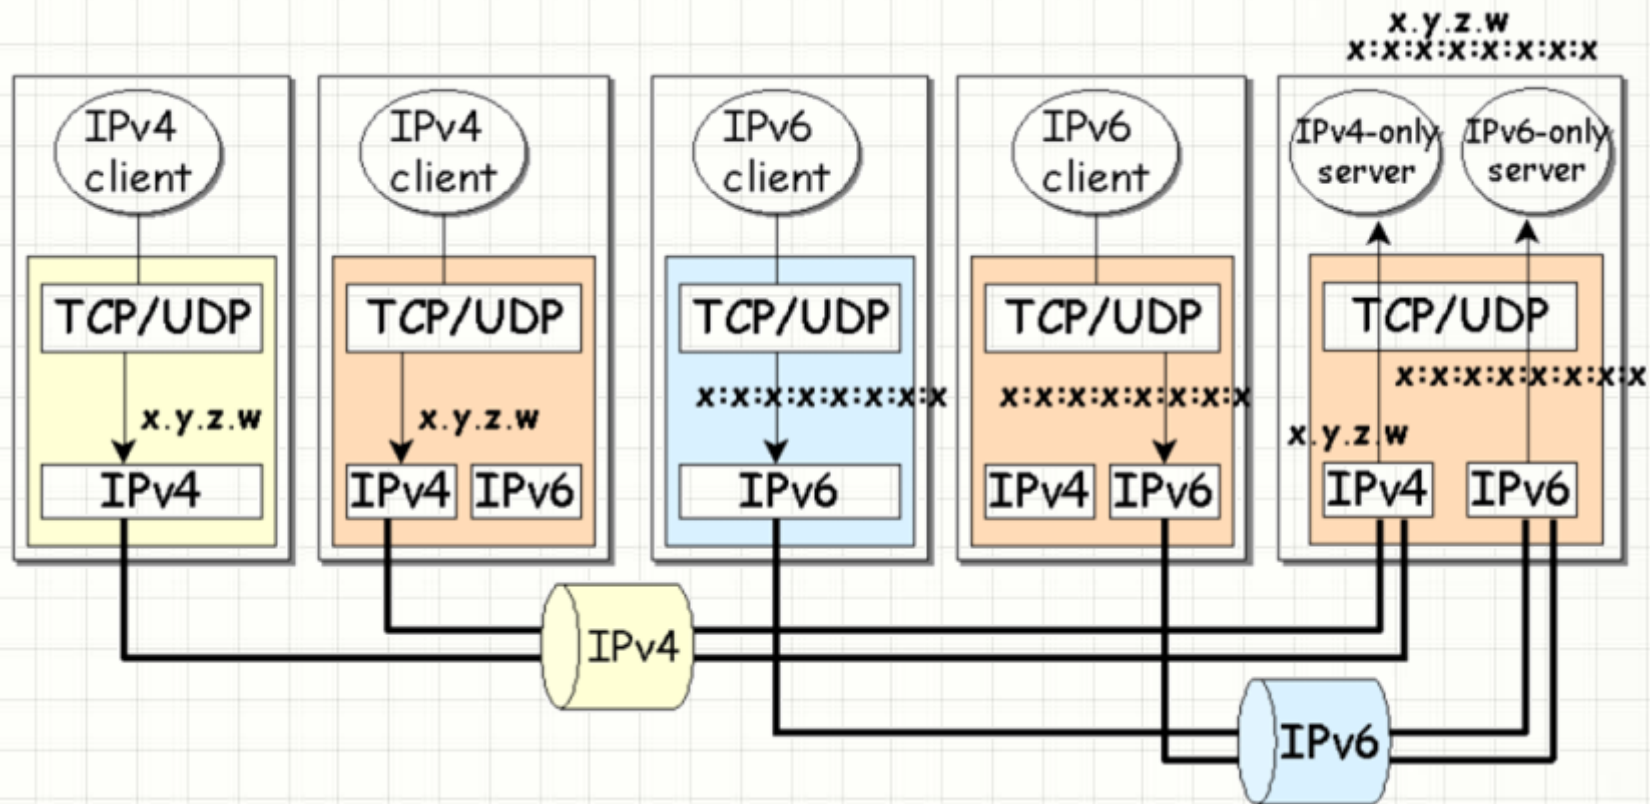
IPv4/IPv6 Interoperability (IPv4 server at dual stack)



IPv4/IPv6 Interoperability (IPv6 server at dual stack)



IPv4/IPv6 Interoperability (IPv4-only server and IPv6-only server at dual stack)



IPv4/IPv6 Interoperability

Summary

		IPv4 server application		IPv6 server application	
		IPv4 node	Dual-stack	IPv6 node	Dual-stack
IPv4 client	IPv4 node	IPv4	IPv4	X	IPv4
	Dual-stack	IPv4	IPv4	X	IPv4
IPv6 client	IPv6 node	X	X	IPv6	IPv6
	Dual-stack	IPv4	IPv4 / X	IPv6	IPv6



Source Code Porting

Socket Address Structures (IPv6)

```
struct in6_addr {
    union {
        uint8_t u6_addr8[16];
        uint16_t u6_addr16[8];
        uint32_t u6_addr32[4];
    } in6_u;

#define s6_addr                in6_u.u6_addr8
#define s6_addr16              in6_u.u6_addr16
#define s6_addr32              in6_u.u6_addr32
};

struct sockaddr_in6 {
    sa_family_t sin6_family;    /* AF_INET6 */
    in_port_t sin6_port;        /* Transport layer port # */
    uint32_t sin6_flowinfo;     /* IPv6 flow information */
    struct in6_addr sin6_addr;  /* IPv6 address */
    uint32_t sin6_scope_id;     /* IPv6 scope-id */
};
```

Socket Address Structures (Portable address)

```
/* Structure large enough to hold any socket address (with the
historical exception of AF_UNIX). 128 bytes reserved. */

#if ULONG_MAX > 0xffffffff
# define __ss_aligntype __uint64_t
#else
# define __ss_aligntype __uint32_t
#endif
#define _SS_SIZE          128
#define _SS_PADSIZE      (_SS_SIZE - (2 * sizeof (__ss_aligntype)))

struct sockaddr_storage
{
    sa_family_t ss_family;      /* Address family */
    __ss_aligntype __ss_align; /* Force desired alignment. */
    char __ss_padding[_SS_PADSIZE];
};
```

Socket API – no change

```
int      socket   (int domain, int type, int protocol);

int      listen   (int s, int backlog);

ssize_t write     (int fd, const void *buf, size_t count);

int      send     (int s, const void *msg, size_t len, int flags);

int      sendmsg  (int s, const struct msghdr *msg, int flags);

ssize_t read      (int fd, void *buf, size_t count);

int      recv     (int s, void *buf, size_t len, int flags);

int      recvmsg  (int s, struct msghdr *msg, int flags);

int      close    (int fd);

int      shutdown(int s, int how);
```


Socket API – change

- The socket address provided by application to kernel

```
int bind    (int sockfd, struct sockaddr *my_addr, socklen_t addrlen);  
  
int connect(int sockfd, const struct sockaddr *serv_addr,  
            socklen_t addrlen);  
  
int sendto (int s, const void *msg, size_t len, int flags,  
            const struct sockaddr *to, socklen_t tolen);
```

- The socket address provided by kernel to application

```
int accept  (int s, struct sockaddr *addr, socklen_t *addrlen);  
  
int recvfrom (int s, void *buf, size_t len, int flags,  
             struct sockaddr *from, socklen_t *fromlen);  
  
int getpeername(int s, struct sockaddr *name, socklen_t *namelen);  
  
int getsockname(int s, struct sockaddr *name, socklen_t *namelen);
```

Typical update when porting to IPv6

➤ int **socket** (int domain, int type, int protocol);

- IPv4 source code:

```
socket(PF_INET, SOCK_STREAM, 0); /* TCP socket */  
socket(PF_INET, SOCK_DGRAM, 0); /* UDP socket */
```

- IPv6 source code:

```
socket(PF_INET6, SOCK_STREAM, 0); /* TCP socket */  
socket(PF_INET6, SOCK_DGRAM, 0); /* UDP socket */
```

Typical update when porting to IPv6

➤ int **bind** (int sockfd, struct sockaddr *my_addr, socklen_t addrlen);

- IPv4 source code

```
struct sockaddr_in addr;  
socklen_t      addrlen = sizeof(addr);  
  
/*  
    fill addr structure using an IPv4 address before calling socket  
    function  
*/  
  
bind(sockfd, (struct sockaddr *)&addr, addrlen);
```

- IPv6 source code:

```
struct sockaddr_in6 addr;  
socklen_t      addrlen = sizeof(addr);  
  
/*  
    fill addr structure using an IPv6 address before calling socket  
    function  
*/  
  
bind(sockfd, (struct sockaddr *)&addr, addrlen);
```

Typical update when porting to IPv6

➤ int **bind** (int sockfd, struct sockaddr *my_addr, socklen_t addrlen);

- Portable source code

```
struct sockaddr_storage addr;
socklen_t          addrlen= sizeof(addr);

/*
    fill addr structure using an IPv4/IPv6 address and
    fill addrlen before calling socket function
*/

bind(sockfd, (struct sockaddr *)&addr, addrlen);
```

Typical update when porting to IPv6

➤ int **accept** (int s, struct sockaddr *addr, socklen_t *addrlen);

- IPv4 source code

```
struct sockaddr_in addr;  
socklen_t      addrlen = sizeof(addr);  
  
accept(sockfd, (struct sockaddr *)&addr, &addrlen);  
  
/*  
   addr structure contains an IPv4 address  
*/
```

- IPv6 source code:

```
struct sockaddr_in6 addr;  
socklen_t      addrlen = sizeof(addr);  
  
accept(sockfd, (struct sockaddr *)&addr, &addrlen);  
  
/*  
   addr structure contains an IPv4 address  
*/
```


Typical update when porting to IPv6

➤ int **accept** (int s, struct sockaddr *addr, socklen_t *addrlen);

- Portable source code

```
struct sockaddr_storage addr;  
socklen_t                addrlen = sizeof(addr);  
  
accept(sockfd, (struct sockaddr *)&addr, &addrlen);  
  
/*  
   addr structure contains an IPv4/IPv6 address  
   addrlen contains the size of the addr structure returned  
*/
```

Address conversion functions

➤ IPv4 address conversion functions

```
/*
    From text to IPv4 binary representation
*/
int      inet_aton (const char *cp, struct in_addr *inp);
in_addr_t inet_addr( const char *cp);

/*
    From IPv4 binary to text representation
*/
char      *inet_ntoa(struct in_addr in);
```

➤ IPv4 /IPv6 address conversion functions

```
/*
    From presentation to IPv4/IPv6 binary representation
*/
int inet_pton(int family, const char *src, void *dst);

/*
    From IPv4/IPv6 binary to presentation
*/
const char *inet_ntop(int family, const void *src,
                      char *dst, size_t cnt);
```

Address conversion functions

- IPv4 source code

```
struct sockaddr_in addr;  
char *straddr;  
  
memset(&addr, 0, sizeof(addr));  
addr.sin_family = AF_INET;      /* family */  
addr.sin_port = htons(MYPORT);  /* port, network byte order */  
  
/*  
    from text to binary representation  
*/  
inet_aton("138.4.2.10", &(addr.sin_addr));  
  
/*  
    from binary to text representation  
*/  
straddr = inet_ntoa(addr.sin_addr);
```

Address conversion functions

- IPv6 source code:

```
struct sockaddr_in6 addr;  
char straddr[INET6_ADDRSTRLEN];  
  
memset(&addr, 0, sizeof(addr));  
addr.sin6_family = AF_INET6;      /* family */  
addr.sin6_port = htons(MYPORT);   /* port, network byte order */  
  
/*  
    from presentation to binary representation  
*/  
inet_pton(AF_INET6, "2001:720:1500:1::a100",  
          &(addr.sin6_addr));  
  
/*  
    from binary representation to presentation  
*/  
inet_ntop(AF_INET6, &addr.sin6_addr, straddr,  
          sizeof(straddr));
```

Name resolving

```
struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);
```

```
struct hostent *gethostbyname(const char *name);
```



```
struct addrinfo {
    int      ai_flags;          /* AI_PASSIVE, AI_CANONNAME */
    int      ai_family;         /* AF_UNSPEC, AF_INET, AF_INET6 */
    int      ai_socktype;       /* SOCK_STREAM, SOCK_DGRAM ... */
    int      ai_protocol;       /* IPPROTO_IP, IPPROTO_IPV6 */
    size_t   ai_addrlen;        /* length of ai_addr */
    struct sockaddr ai_addr;     /* socket address structure */
    char     ai_canonname;       /* canonical name */
    struct addrinfo ai_next;     /* next addrinfo structure */
};

/* function to get socket address structures */

int getaddrinfo(const char *node, const char *service,
               const struct addrinfo *hints,
               struct addrinfo **res);

/* function to free the resources allocated by getaddrinfo */

void freeaddrinfo(struct addrinfo *res);
```


Name resolving

```
n = getaddrinfo(hostname, service, &hints, &res);

/*
   Try open socket with each address getaddrinfo returned,
   until getting a valid socket.
*/

resave = res;

while (res) {
    sockfd = socket(res->ai_family,
                    res->ai_socktype,
                    res->ai_protocol);

    if (!(sockfd < 0))
        break;

    res = res->ai_next;
}

freeaddrinfo(ressave);
```

Application Deployment

Application		Deploy Node		Connections	
Network API used	Application type	IPv6 kernel support	Dual stack activated	IPv4 connection	IPv6 connection
IPv6 extensions (portable code)	IPv6-enabled	Yes	Yes	IPv4 stack	IPv6 stack
IPv6 extensions (portable code)	IPv6-enabled		No (only IPv4 is activated)	IPv4 stack	Address resolution successful Connection error
Without IPv6 extensions (old API)	IPv4-only	No		IPv4 stack	Error

Q: How to know whether the linux kernel support ipv6?

A: `ls -lrt /proc/net/ipv6`

Q: How to know whether the ipv6 stack be activated?

A: `ping6 -c5 ::1`


Summary

IPv6 Addressing

Source Code Porting

- Socket Address Structures
- Socket API
- Address Conversion Functions
- Name Resolving
- Application Deployment

Transition &
Interoperability





QUESTIONS?